

Group 7 Project Submission

Student names: JANE MARTHA
JAMES NJORGE
LEON MAINA
DANIEL WAHOME
MAGDALENE ONDIMU
FAITH MWENDA

Student pace: PART TIME

Instructor name: NOAH KANDIE
WILLIAM OKOMBA

import standard packages

```
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
import numpy as np
import warnings
warnings.filterwarnings('ignore')
import matplotlib.pyplot as plt
plt.style.use('ggplot')
from scipy import stats
import statsmodels.api as sm
import statsmodels.graphics as smg
import statsmodels.stats.api as sms
from statsmodels.stats.outliers_influence import
variance_inflation_factor
import datetime as dt
```

#Loading data set for Analysis

```
kc = pd.read_csv(r"C:\Users\HP\Documents\Flatiron\Assignments\Phase 2
Project\kc_house_data.csv")
```

```
kc.head(5)
```

	id	date	price	bedrooms	bathrooms	sqft_living
sqft_lot \						
0	7129300520	10/13/2014	221900	3	1.00	1180
5650						
1	6414100192	12/9/2014	538000	3	2.25	2570
7242						
2	5631500400	2/25/2015	180000	2	1.00	770
10000						
3	2487200875	12/9/2014	604000	4	3.00	1960
5000						
4	1954400510	2/18/2015	510000	3	2.00	1680

8080

	floors	waterfront	view	...	grade	sqft_above
sqft_basement \						
0	1.0	NaN	NONE	...	7 Average	1180
0						
1	2.0	NO	NONE	...	7 Average	2170
400						
2	1.0	NO	NONE	...	6 Low Average	770
0						
3	1.0	NO	NONE	...	7 Average	1050
910						
4	1.0	NO	NONE	...	8 Good	1680
0						

	yr_built	yr_renovated	zipcode	lat	long	sqft_living15
sqft_lot15						
0	1955	0.0	98178	47.5112	-122.257	1340
5650						
1	1951	1991.0	98125	47.7210	-122.319	1690
7639						
2	1933	NaN	98028	47.7379	-122.233	2720
8062						
3	1965	0.0	98136	47.5208	-122.393	1360
5000						
4	1987	0.0	98074	47.6168	-122.045	1800
7503						

[5 rows x 21 columns]

#Data understanding and exploration

kc.info()

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 21597 entries, 0 to 21596

Data columns (total 21 columns):

#	Column	Non-Null Count	Dtype
0	id	21597 non-null	int64
1	date	21597 non-null	object
2	price	21597 non-null	int64
3	bedrooms	21597 non-null	int64
4	bathrooms	21597 non-null	float64
5	sqft_living	21597 non-null	int64
6	sqft_lot	21597 non-null	int64
7	floors	21597 non-null	float64
8	waterfront	19221 non-null	object
9	view	21534 non-null	object
10	condition	21597 non-null	object
11	grade	21597 non-null	object

```

12  sqft_above      21597 non-null  int64
13  sqft_basement   21597 non-null  object
14  yr_built        21597 non-null  int64
15  yr_renovated     17755 non-null  float64
16  zipcode         21597 non-null  int64
17  lat            21597 non-null  float64
18  long           21597 non-null  float64
19  sqft_living15    21597 non-null  int64
20  sqft_lot15       21597 non-null  int64

```

dtypes: float64(5), int64(10), object(6)

memory usage: 3.5+ MB

kc.describe()

	id	price	bedrooms	bathrooms
count	2.159700e+04	2.159700e+04	21597.000000	21597.000000
mean	4.580474e+09	5.402966e+05	3.373200	2.115826
std	2.876736e+09	3.673681e+05	0.926299	0.768984
min	1.000102e+06	7.800000e+04	1.000000	0.500000
25%	2.123049e+09	3.220000e+05	3.000000	1.750000
50%	3.904930e+09	4.500000e+05	3.000000	2.250000
75%	7.308900e+09	6.450000e+05	4.000000	2.500000
max	9.900000e+09	7.700000e+06	33.000000	8.000000

	sqft_lot	floors	sqft_above	yr_built
count	2.159700e+04	21597.000000	21597.000000	21597.000000
mean	1.509941e+04	1.494096	1788.596842	1970.999676
std	4.141264e+04	0.539683	827.759761	29.375234
min	5.200000e+02	1.000000	370.000000	1900.000000
25%	5.040000e+03	1.000000	1190.000000	1951.000000
50%	7.618000e+03	1.500000	1560.000000	1975.000000
75%	1.068500e+04	2.000000	2210.000000	1997.000000
max	1.651359e+06	3.500000	9410.000000	2015.000000

2015.000000

	zipcode	lat	long	sqft_living15
sqft_lot15				
count	21597.000000	21597.000000	21597.000000	21597.000000
mean	98077.951845	47.560093	-122.213982	1986.620318
std	53.513072	0.138552	0.140724	685.230472
min	98001.000000	47.155900	-122.519000	399.000000
25%	98033.000000	47.471100	-122.328000	1490.000000
50%	98065.000000	47.571800	-122.231000	1840.000000
75%	98118.000000	47.678000	-122.125000	2360.000000
max	98199.000000	47.777600	-121.315000	6210.000000

#changing the selling date to and updating the column name to yr_sold

```
kc['date']=pd.to_datetime(kc['date'])
```

```
kc['date'] = kc['date'].dt.year
```

```
kc.rename (columns={'date': 'yr_sold'}, inplace=True)
```

```
kc['yr_sold']= kc['yr_sold'].astype(int)
```

```
kc.head(5)
```

	id	yr_sold	price	bedrooms	bathrooms	sqft_living
sqft_lot \						
0	7129300520	2014	221900	3	1.00	1180
1	6414100192	2014	538000	3	2.25	2570
2	5631500400	2015	180000	2	1.00	770
3	2487200875	2014	604000	4	3.00	1960
4	1954400510	2015	510000	3	2.00	1680

	floors	waterfront	view	...	grade	sqft_above
sqft_basement \						
0	1.0	NaN	NONE	...	7 Average	1180
1	2.0	NO	NONE	...	7 Average	2170
2	1.0	NO	NONE	...	6 Low Average	770

```

0
3      1.0          NO  NONE  ...      7 Average      1050
910
4      1.0          NO  NONE  ...      8 Good      1680
0

   yr_built  yr_renovated  zipcode      lat      long  sqft_living15
sqft_lot15
0    1955          0.0    98178  47.5112 -122.257          1340
5650
1    1951        1991.0    98125  47.7210 -122.319          1690
7639
2    1933          NaN    98028  47.7379 -122.233          2720
8062
3    1965          0.0    98136  47.5208 -122.393          1360
5000
4    1987          0.0    98074  47.6168 -122.045          1800
7503

```

[5 rows x 21 columns]

```
kc.isna().sum()
```

```

id                0
yr_sold           0
price             0
bedrooms          0
bathrooms         0
sqft_living       0
sqft_lot          0
floors            0
waterfront       2376
view              63
condition         0
grade             0
sqft_above        0
sqft_basement     0
yr_built          0
yr_renovated     3842
zipcode           0
lat               0
long              0
sqft_living15     0
sqft_lot15        0
dtype: int64

```

```

#Filling missing values in the 'view' and 'waterfront' columns
kc['waterfront'].fillna('NO', inplace=True)
kc['view'].fillna('NONE', inplace=True)
kc.isna().sum()

```

```

id                0
yr_sold           0
price             0
bedrooms          0
bathrooms         0
sqft_living       0
sqft_lot          0
floors            0
waterfront        0
view              0
condition         0
grade             0
sqft_above        0
sqft_basement     0
yr_built          0
yr_renovated      3842
zipcode           0
lat               0
long              0
sqft_living15     0
sqft_lot15        0
dtype: int64

```

#Counting the occurrences of unique values in 'yr_renovated'

```
kc['yr_renovated'].value_counts()
```

```

0.0      17011
2014.0     73
2003.0     31
2013.0     31
2007.0     30
...
1946.0      1
1959.0      1
1971.0      1
1951.0      1
1954.0      1

```

```
Name: yr_renovated, Length: 70, dtype: int64
```

```
kc['yr_renovated'].fillna(0.0, inplace=True)
```

```
kc['yr_renovated'].astype(int)
```

```

0         0
1      1991
2         0
3         0
4         0
...
21592     0

```

```
21593      0
21594      0
21595      0
21596      0
Name: yr_renovated, Length: 21597, dtype: int32
```

```
kc.isna().sum()
```

```
id          0
yr_sold     0
price       0
bedrooms    0
bathrooms   0
sqft_living 0
sqft_lot    0
floors      0
waterfront  0
view        0
condition   0
grade       0
sqft_above  0
sqft_basement 0
yr_built    0
yr_renovated 0
zipcode     0
lat         0
long        0
sqft_living15 0
sqft_lot15  0
dtype: int64
```

```
yr_renovated = kc['yr_renovated']
kc['renovated_last_10'] = (yr_renovated >= (kc['yr_sold'] - 10))
kc['renovated_last_10'] = kc['renovated_last_10'].map({True: 'Yes',
False: 'No'})
```

```
kc.head(5)
```

	id	yr_sold	price	bedrooms	bathrooms	sqft_living
0	7129300520	2014	221900	3	1.00	1180
1	6414100192	2014	538000	3	2.25	2570
2	5631500400	2015	180000	2	1.00	770
3	2487200875	2014	604000	4	3.00	1960
4	1954400510	2015	510000	3	2.00	1680

	floors	waterfront	view	...	sqft_above	sqft_basement	yr_built	\
0	1.0	NO	NONE	...	1180	0	1955	
1	2.0	NO	NONE	...	2170	400	1951	
2	1.0	NO	NONE	...	770	0	1933	
3	1.0	NO	NONE	...	1050	910	1965	
4	1.0	NO	NONE	...	1680	0	1987	

	yr_renovated	zipcode	lat	long	sqft_living15	
0	0.0	98178	47.5112	-122.257	1340	5650
1	1991.0	98125	47.7210	-122.319	1690	7639
2	0.0	98028	47.7379	-122.233	2720	8062
3	0.0	98136	47.5208	-122.393	1360	5000
4	0.0	98074	47.6168	-122.045	1800	7503

	renovated_last_10
0	No
1	No
2	No
3	No
4	No

[5 rows x 22 columns]

```
kc['Age'] = kc['yr_sold']-kc['yr_built']
```

```
kc.head(5)
```

	id	yr_sold	price	bedrooms	bathrooms	sqft_living
0	7129300520	2014	221900	3	1.00	1180
1	6414100192	2014	538000	3	2.25	2570
2	5631500400	2015	180000	2	1.00	770
3	2487200875	2014	604000	4	3.00	1960
4	1954400510	2015	510000	3	2.00	1680

	floors	waterfront	view	...	sqft_basement	yr_built	yr_renovated
0	1.0	NO	NONE	...	0	1955	0.0

1	2.0	NO	NONE	...	400	1951	1991.0
98125							
2	1.0	NO	NONE	...	0	1933	0.0
98028							
3	1.0	NO	NONE	...	910	1965	0.0
98136							
4	1.0	NO	NONE	...	0	1987	0.0
98074							

	lat	long	sqft_living15	sqft_lot15	renovated_last_10	Age
0	47.5112	-122.257	1340	5650	No	59
1	47.7210	-122.319	1690	7639	No	63
2	47.7379	-122.233	2720	8062	No	82
3	47.5208	-122.393	1360	5000	No	49
4	47.6168	-122.045	1800	7503	No	28

[5 rows x 23 columns]

#Dropping columns that are not needed

columns_to_drop =

`['id', 'zipcode', 'lat', 'long', 'sqft_living15', 'sqft_lot15', 'sqft_basement', 'yr_sold', 'yr_renovated']`

`kc.drop(columns_to_drop, axis = 1, inplace = True)`

#Confirming the new dataframe

`kc.head(5)`

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors
waterfront \						
0	221900	3	1.00	1180	5650	1.0
NO						
1	538000	3	2.25	2570	7242	2.0
NO						
2	180000	2	1.00	770	10000	1.0
NO						
3	604000	4	3.00	1960	5000	1.0
NO						
4	510000	3	2.00	1680	8080	1.0
NO						

	view	condition	grade	sqft_above	yr_built
renovated_last_10					
Age					
0	NONE	Average	7	Average	1180
No	59				1955
1	NONE	Average	7	Average	2170
					1951


No	63						
2	NONE	Average	6	Low	Average	770	1933
No	82						
3	NONE	Very Good	7	Average		1050	1965
No	49						
4	NONE	Average	8	Good		1680	1987
No	28						

Visualizing

```
#Creating a plotting function for ease or resizing in different plots
def resizeplot(l,a):
    plt.figure(figsize=(l,a));

#Creating a plot that will visualize the relationship between price
and sqft_lot
resizeplot(6,4)
sns.relplot(x='sqft_lot',y='price',data=kc,palette='terrain');
plt.show()

<Figure size 432x288 with 0 Axes>
```



```
#Finding out the occurrences of bedroom values
bedrooms_counts = kc['bedrooms'].value_counts()
bedrooms_counts
```

```
3      9824
4      6882
2      2760
5      1601
6       272
1       196
7        38
8         13
9          6
10         3
11         1
33         1
```

```
Name: bedrooms, dtype: int64
```

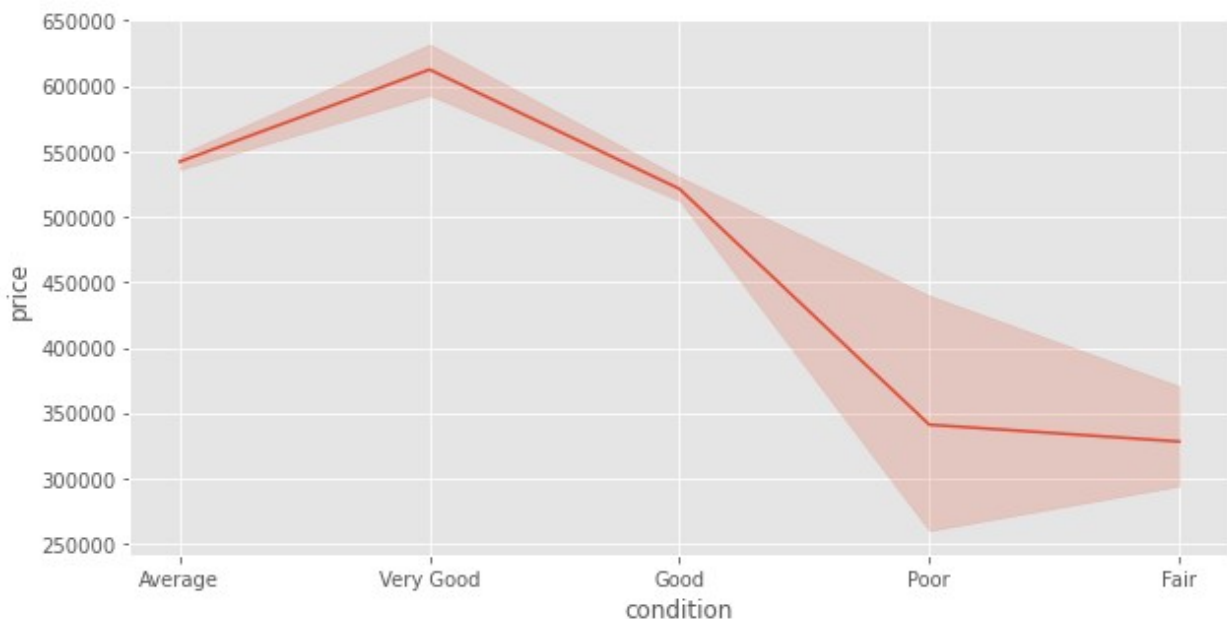
```
#Confirming how the conditon of the house over the year affects the price
```

```
#We can see that if the conditon of the house- is improved then the price of the house is higher
```

```
#Poor conditons leads to lower price
```


```
resizeplot(10,5)
```

```
sns.lineplot(x='condition',y='price',data=kc,palette='terrain')
plt.show()
```




```
#Using histogram to visualize the distribution of the price
#As the count is high then the price increases
```

```
resizeplot(6,4)
sns.histplot(kc['price'],kde=True,bins=50, color = 'red');
```




*#Using linear to analyse price increament over the year
#At the beggining of the yrs the price is low as the yrs increases
then the price has a higher peak*

```
resizeplot(12,6)
sns.lineplot(x='Age',y='price',data=kc);
```



*#Price trend over the years
 #The price from 2000 to 2015 was good meaning the owners of real
 estate were getting good profit
 #It seems that the houses had been reinovated*

```
resizeplot(12,6)
sns.lineplot(x='yr_built',y='price',data=kc,palette = 'deep');
```




*#Showing the price changes in relation to unit grade
 # poor grade causes the price to decrease and vice versa*

```

resizeplot(12,4)# changed the size to give a better view of the grade.
sns.lineplot(x='grade',y='price',data=kc,palette='terrain');
plt.xticks(rotation=45)
plt.show()

```



Categorical Variables

Plotting histogram for columns within the dataset
#The histogram shows us how different features in each column affect each other

```

fig, axes = plt.subplots(nrows=(5), ncols=3, figsize=(20,10))
df_cols = kc.columns

```

Using sns color pallets for each plot

```

color = sns.color_palette("Blues", n_colors=1)[0]

```

creation of a function for plotting the hustogram for the given columns

```

for col, ax in zip(df_cols, axes.flatten()):
    ax.hist(kc[col].dropna(), bins='auto', color=color )
    ax.set_title(col)

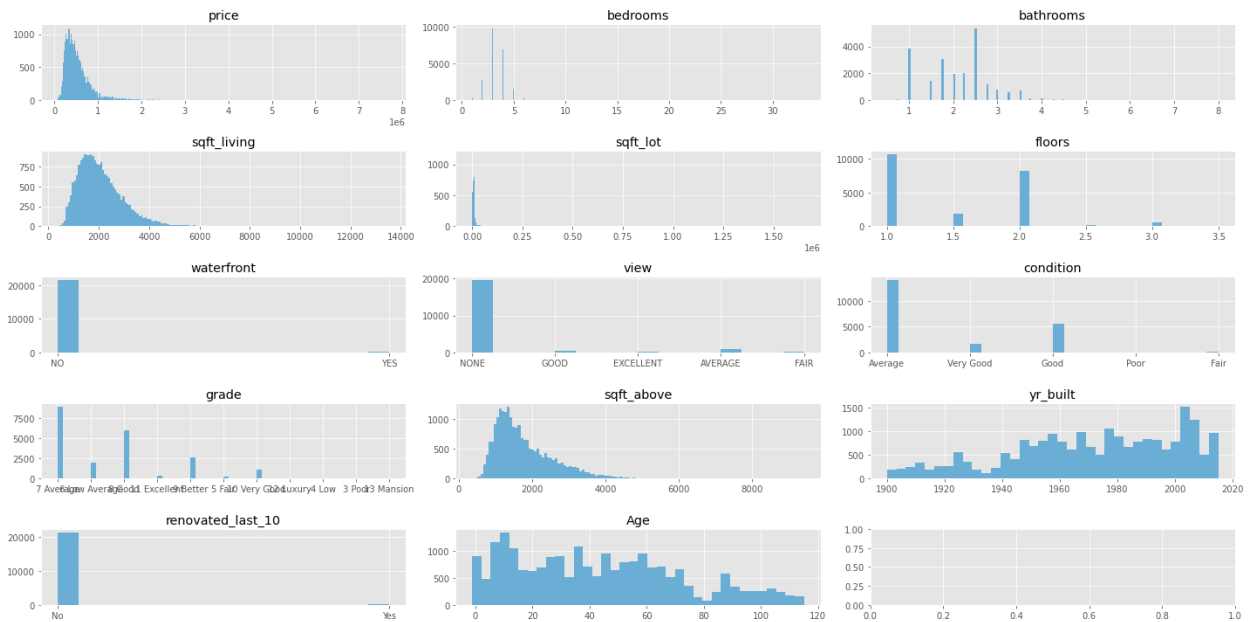
```

automatically adjusting subplot params so that the subplot(s) fits in to the figure area

```

fig.tight_layout()

```



*#From the above plots we can see that
floors,waterfront,condition,grade,renovated_last_10,bedrooms,bathrooms
are categorical*

*#Creating a category of values to work with
categories = ['bedrooms', 'bathrooms', 'floors', 'waterfront', 'view',
'condition', 'grade', 'renovated_last_10']*

for cate in categories:

getting the value counts

counts = kc[cate].value_counts()

Isolate offending categories for each variable

bad_cate = counts[counts < 50].index

*# Isolate indices in the dataset where offending categories are
found*

to_drop = kc[kc[cate].isin(bad_cate)].index

Dropping unnecessary data within the category in dataset

kc.drop(to_drop, inplace=True)

converting view into binary (binarzing)

#for consistent and easier to compare in the view.

*view_dict = {
 'FAIR': 1,
 'AVERAGE': 1,
 'GOOD': 1,*

```

        'EXCELLENT': 1,
        'NONE': 0
    }

    kc['view'] = kc['view'].map(view_dict)

    # Binarizing waterfront
    waterfront_dict = {
        'YES': 1,
        'NO': 0
    }

    kc['waterfront'] = kc['waterfront'].map(waterfront_dict)

    # Binarizing renovated_last_10
    renovated_dict = {
        'Yes': 1,
        'No': 0
    }

    kc['renovated_last_10'] = kc['renovated_last_10'].map(renovated_dict)

```

OneHotEncoding

```

#Creating a copy of the dataset so as not to alter the original copy incase
kc2= kc.copy()

```

Generating dummies

```

#Converting the 3columns: bedrooms,bathrooms,floors to string for pandas to be able to dumify them
col = ['bedrooms', 'bathrooms', 'floors']
kc[col] = kc[col].astype(str)

```

```

#Checking if the conversion was successful
kc.dtypes

```

price	int64
bedrooms	object
bathrooms	object
sqft_living	int64
sqft_lot	int64
floors	object
waterfront	int64
view	int64
condition	object
grade	object
sqft_above	int64
yr_built	int64


```
renovated_last_10    int64
Age                  int64
dtype: object
```

```
# Creating variable from the already cleaned original dataset
kc_binary = kc[['waterfront', 'view', 'renovated_last_10']]
kc_num = kc[['price', 'sqft_living', 'sqft_lot', 'Age']]
kc_cate = kc[['floors', 'bedrooms', 'bathrooms', 'condition',
'grade']]
```

```
# Applying one-hot encoding to the categorical features
kc_cate_dummies = pd.get_dummies(kc_cate, dtype=int)
```

```
#Creating a list of dummies to be dropped
```

```
dummies_to_drop = [
    'floors_1.0',
    'bedrooms_1',
    'bathrooms_0.75',
    'condition_Fair',
    'grade_5 Fair'
]
```

```
#Dropping the specified dummies
```

```
kc_cate_dummies.drop(
    dummies_to_drop,
    axis = 1,
    inplace=True)
```

```
#Combining the variable into a single variable
```

```
kc = pd.concat([kc_num, kc_binary, kc_cate_dummies], axis=1)
```

```
#Confirming if the event was successful
```

```
kc.head()
```

	price	sqft_living	sqft_lot	Age	waterfront	view
renovated_last_10 \						
0	221900	1180	5650	59	0	0
0						
1	538000	2570	7242	63	0	0
0						
2	180000	770	10000	82	0	0
0						
3	604000	1960	5000	49	0	0
0						
4	510000	1680	8080	28	0	0
0						

	floors_1.5	floors_2.0	floors_2.5	...	condition_Average
condition_Good \					
0	0	0	0	...	1

```

0
1      0      1      0 ...      1
0
2      0      0      0 ...      1
0
3      0      0      0 ...      0
0
4      0      0      0 ...      1
0

    condition_Very Good    grade_10 Very Good    grade_11 Excellent \
0              0              0              0
1              0              0              0
2              0              0              0
3              1              0              0
4              0              0              0

    grade_12 Luxury    grade_6 Low Average    grade_7 Average    grade_8 Good
\
0              0              0              1              0
1              0              0              1              0
2              0              1              0              0
3              0              0              1              0
4              0              0              0              1

    grade_9 Better
0              0
1              0
2              0
3              0
4              0

[5 rows x 40 columns]

#Confirming the columns
kc.columns

Index(['price', 'sqft_living', 'sqft_lot', 'Age', 'waterfront',
      'view',
      'renovated_last_10', 'floors_1.5', 'floors_2.0', 'floors_2.5',
      'floors_3.0', 'bedrooms_2', 'bedrooms_3', 'bedrooms_4',
      'bedrooms_5',
      'bedrooms_6', 'bathrooms_1.0', 'bathrooms_1.5',
      'bathrooms_1.75',
      'bathrooms_2.0', 'bathrooms_2.25', 'bathrooms_2.5',
      'bathrooms_2.75',

```

```

        'bathrooms_3.0', 'bathrooms_3.25', 'bathrooms_3.5',
        'bathrooms_3.75',
        'bathrooms_4.0', 'bathrooms_4.25', 'bathrooms_4.5',
        'condition_Average',
        'condition_Good', 'condition_Very Good', 'grade_10 Very Good',
        'grade_11 Excellent', 'grade_12 Luxury', 'grade_6 Low Average',
        'grade_7 Average', 'grade_8 Good', 'grade_9 Better'],
        dtype='object')

#Setting the columns in an ascending order for easy analysis
grade_columns = [
    'grade_6 Low Average',
    'grade_7 Average',
    'grade_8 Good',
    'grade_9 Better',
    'grade_10 Very Good',
    'grade_11 Excellent',
    'grade_12 Luxury'
]

# Extracting other columns not related to 'grade'

other_columns = [col for col in kc.columns if col not in
grade_columns]

# Reordering columns

reordered_columns = other_columns + grade_columns
kc = kc[reordered_columns]

#Accessing the columns
kc.columns

Index(['price', 'sqft_living', 'sqft_lot', 'Age', 'waterfront',
'view',
      'renovated_last_10', 'floors_1.5', 'floors_2.0', 'floors_2.5',
      'floors_3.0', 'bedrooms_2', 'bedrooms_3', 'bedrooms_4',
'bedrooms_5',
      'bedrooms_6', 'bathrooms_1.0', 'bathrooms_1.5',
'bedrooms_1.75',
      'bathrooms_2.0', 'bathrooms_2.25', 'bathrooms_2.5',
'bedrooms_2.75',
      'bathrooms_3.0', 'bathrooms_3.25', 'bathrooms_3.5',
'bedrooms_3.75',
      'bathrooms_4.0', 'bathrooms_4.25', 'bathrooms_4.5',
'condition_Average',
      'condition_Good', 'condition_Very Good', 'grade_6 Low Average',
      'grade_7 Average', 'grade_8 Good', 'grade_9 Better',
      'grade_10 Very Good', 'grade_11 Excellent', 'grade_12 Luxury'],
      dtype='object')

```

#Accessing the columns and rows

kc.head()

	price	sqft_living	sqft_lot	Age	waterfront	view
renovated_last_10 \						
0	221900	1180	5650	59	0	0
0						
1	538000	2570	7242	63	0	0
0						
2	180000	770	10000	82	0	0
0						
3	604000	1960	5000	49	0	0
0						
4	510000	1680	8080	28	0	0
0						

	floors_1.5	floors_2.0	floors_2.5	...	condition_Average
condition_Good \					
0	0	0	0	...	1
0					
1	0	1	0	...	1
0					
2	0	0	0	...	1
0					
3	0	0	0	...	0
0					
4	0	0	0	...	1
0					

	condition_Very Good	grade_6 Low Average	grade_7 Average	grade_8
Good \				
0	0	0	1	
0				
1	0	0	1	
0				
2	0	1	0	
0				
3	1	0	1	
0				
4	0	0	0	
1				

	grade_9 Better	grade_10 Very Good	grade_11 Excellent	grade_12
Luxury				
0	0	0	0	
0				
1	0	0	0	
0				
2	0	0	0	
0				

```

3          0          0          0
0
4          0          0          0
0

```

[5 rows x 40 columns]

Modeling


#Confirming the linear relationship between variable and price

```
kc_num.head()
```

	price	sqft_living	sqft_lot	Age
0	221900	1180	5650	59
1	538000	2570	7242	63
2	180000	770	10000	82
3	604000	1960	5000	49
4	510000	1680	8080	28

#Plotting scatter plot for clear visulaization of different features

```
pd.plotting.scatter_matrix(kc_num, figsize=(20,15), alpha=.3);
```



#Displaying the correlation matrix in response to price in descending order


```
kc_num.corr()['price'].sort_values(ascending=False)
```

```
price      1.000000
sqft_living 0.682342
sqft_lot    0.084739
Age        -0.048693
Name: price, dtype: float64
```

#Ensuring proper visualization using heatmap
#the lighter the color the stronger the correlation

```
fig, ax = plt.subplots(figsize = (8,10))
```

```
sns.heatmap(
    kc_num.corr().abs(),
    # mask=np.triu(np.ones_like(data.corr()), dtype=bool)),
    ax=ax,
    annot=True,
    cbar_kws={"label": "Correlation", "orientation": "horizontal",
"pad": .2, "extend": "both"}
);
```



**sqft_living: 0.682342 (Strong positive correlation) sqft_lot: 0.084739 (Weak positive correlation) Age: -0.048722 (Weak negative correlation)

Assessing Multicollinearity Across Predictor Combinations

```
# Create the correlation matrix directly and then reshape it for
# visualization
# Dropping the 'price' column
df = kc_num.drop('price', axis=1)
```

```

# Creating a correlation matrix
corr_matrix = df.corr().abs()

# Reshaping the correlation matrix for visualization
df_predictor = corr_matrix.stack().reset_index()
df_predictor.columns = ['Variable 1', 'Variable 2', 'Coefficient']

# Dropping duplicate rows where variables are the same
df_predictor = df_predictor[df_predictor['Variable 1'] !=
df_predictor['Variable 2']]

# Sorting by coefficient in descending order
df_predictor.sort_values(by='Coefficient', ascending=False,
inplace=True)

df_predictor.head()

```

	Variable 1	Variable 2	Coefficient
2	sqft_living	Age	0.326558
6	Age	sqft_living	0.326558
1	sqft_living	sqft_lot	0.164552
3	sqft_lot	sqft_living	0.164552
5	sqft_lot	Age	0.049919

```

#sqft_lot and Age lack linear relationship
kc_num.drop(['sqft_lot', 'Age'], axis=1, inplace=True)

def reg_qq_sced(y, X, add_constant=True, qq=True, sced=True):
    """
    Fits a linear regression model, display its summary, and output
    plots to check linear regression assumptions.

    Parameters:
    - y: Target variable.
    - X: Predictor variables.
    - add_constant: Whether to add a constant term to the predictors
    (default: True).
    - qq: Whether to display a QQ plot for residual normality check
    (default: True).
    - sced: Whether to display a plot of predicted values vs.
    residuals for homoscedasticity check (default: True).
    """
    # Add a constant to the predictors if required
    X_sm = sm.add_constant(X, has_constant='add') if add_constant else
X

    # Run a linear regression and display the summary
    model = sm.OLS(y, X_sm).fit()
    display(print(model.summary()))

```



```

# Display a QQ plot for residual normality check
if qq:
    sm.qqplot(model.resid, line='45', fit=True)
    plt.title('QQ plot for residual normality check')
    plt.show()
else:
    pass

# Display a plot of predicted values vs. residuals for
homoscedasticity check
if sced:
    preds = model.predict(X_sm)
    residuals = model.resid
    fig_resid, ax = plt.subplots(figsize=(10, 5))
    fig_resid.suptitle('Predicted vs. residual plot for
homoscedasticity check')
    ax.scatter(preds, residuals, alpha=0.2, color= "blue")
    ax.plot(preds, [0 for _ in range(len(X_sm))])
    ax.set_xlabel("Predicted Value")
    ax.set_ylabel("Actual - Predicted Value")
else:
    pass

# Output additional model performance metrics
print(f'Model adjusted R-squared: {model.rsquared_adj}')
print(f'Model RMSE: {np.sqrt(model.mse_resid)}')

# Set baseline predictor as 'sqft_living'

baseline = 'sqft_living'

# Define target variable and predictor

y = kc.price
X = kc[baseline]

# Feed these inputs into our function

reg_qq_sced(y, X)

```

OLS Regression Results

```

=====
=====
Dep. Variable:                price    R-squared:
0.466
Model:                        OLS      Adj. R-squared:
0.466
Method:                       Least Squares    F-statistic:
1.862e+04

```

Date: Sat, 06 Apr 2024 Prob (F-statistic): 0.00
Time: 11:34:37 Log-Likelihood: -2.9572e+05
No. Observations: 21378 AIC: 5.914e+05
Df Residuals: 21376 BIC: 5.915e+05
Df Model: 1

Covariance Type: nonrobust

=====

=====

	coef	std err	t	P> t	[0.025
--	------	---------	---	------	--------

0.975]

const	-9704.3054	4318.156	-2.247	0.025	-1.82e+04
-1240.397					
sqft_living	262.8635	1.926	136.467	0.000	259.088
266.639					

=====

=====

Omnibus: 12569.130 Durbin-Watson: 1.978
Prob(Omnibus): 0.000 Jarque-Bera (JB): 271238.797
Skew: 2.414 Prob(JB): 0.00
Kurtosis: 19.769 Cond. No. 5.75e+03


=====

Notes:


[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.75e+03. This might indicate that there are strong multicollinearity or other numerical problems.

None



Model adjusted R-squared: 0.4655655910250339
Model RMSE: 246222.6375951674



The linear regression results indicate that the model's R-squared value is 0.466, suggesting that approximately 46.6% of the variance in the target variable (price) is explained by the predictor variable (sqft_living).

The coefficient for the constant term (intercept) is -9704.3054, indicating the predicted price when sqft_living is zero. The coefficient for sqft_living is 262.8635, indicating that for every unit increase in sqft_living, the price is expected to increase by approximately \$262.86. The p-value for sqft_living is less than 0.05, indicating that the predictor variable is statistically significant. The confidence interval for the coefficient of sqft_living ranges from 259.088 to 266.639. The residual plots indicate that there might be some heteroscedasticity present, as the spread of residuals increases with predicted values. The QQ plot suggests that the residuals are approximately normally distributed, but there might be some deviations, especially in the tails.

Overall, the model performs reasonably well, but there might be room for improvement, especially in addressing heteroscedasticity.

```
# Isolating columns to be transformed
```

```
log_trans_cols = ['price', 'sqft_living']  
kc_logged = kc.copy()[log_trans_cols]
```

```
# Log transforming and renaming columns
```

```
kc_logged = np.log(kc_logged)  
kc_logged.columns = kc_logged.columns.map(lambda x: 'log_' + x)
```

```
# Merge it with the rest of the dataset
```

```
kc_transformed = kc_logged.join(kc.drop(log_trans_cols, axis=1))
```

```
kc_transformed.head()
```

	log_price	log_sqft_living	sqft_lot	Age	waterfront	view	\
0	12.309982	7.073270	5650	59	0	0	
1	13.195614	7.851661	7242	63	0	0	
2	12.100712	6.646391	10000	82	0	0	
3	13.311329	7.580700	5000	49	0	0	
4	13.142166	7.426549	8080	28	0	0	

	renovated_last_10	floors_1.5	floors_2.0	floors_2.5	...	\
0	0	0	0	0	...	
1	0	0	1	0	...	
2	0	0	0	0	...	
3	0	0	0	0	...	
4	0	0	0	0	...	

	condition_Average	condition_Good	condition_Very Good	\
0	1	0	0	
1	1	0	0	
2	1	0	0	

3	0	0	1
4	1	0	0

	grade_6 Low Average	grade_7 Average	grade_8 Good	grade_9 Better
\				
0	0	1	0	0
1	0	1	0	0
2	1	0	0	0
3	0	1	0	0
4	0	0	1	0

	grade_10 Very Good	grade_11 Excellent	grade_12 Luxury
0	0	0	0
1	0	0	0
2	0	0	0
3	0	0	0
4	0	0	0

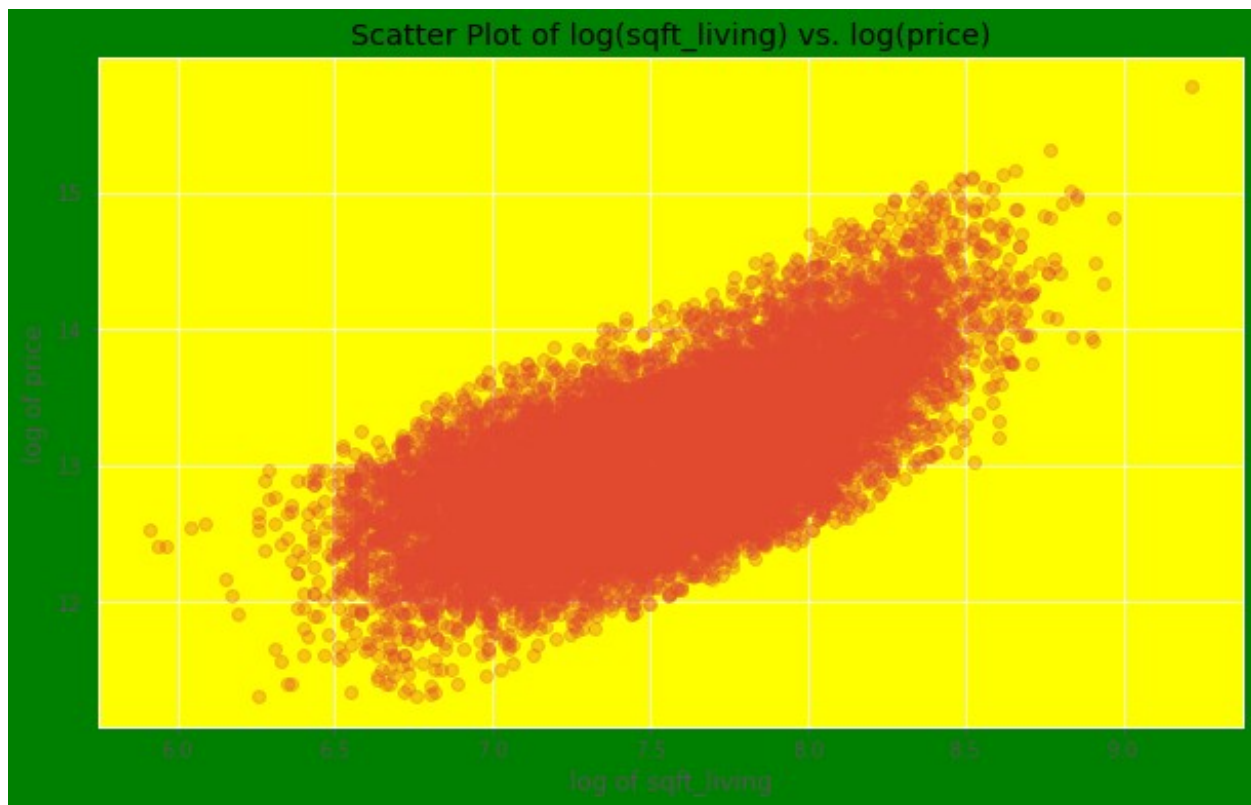
[5 rows x 40 columns]

visualizing linearity between the transformed predictor and target variable
The plot helps visualize the linearity between these two variables, which is essential for linear regression modeling.

```
fig, ax = plt.subplots(figsize= (10,6))
ax = plt.gca()
plt.plot(kc_transformed['log_sqft_living'],
kc_transformed['log_price'], 'o', alpha=0.3)
```

Set the background color of the axis

```
ax.set_facecolor('yellow')
plt.gcf().set_facecolor('green')
plt.xlabel('log of sqft_living')
plt.ylabel('log of price')
plt.title('Scatter Plot of log(sqft_living) vs. log(price)')
plt.show()
```



```
baseline = 'log_sqft_living'

y = kc_transformed.log_price
X = kc_transformed.log_sqft_living
```

Feeding these inputs into our function

```
model = reg_qq_sced(y, X)
```

OLS Regression Results

```
=====
=====
Dep. Variable:          log_price    R-squared:
0.441
Model:                  OLS         Adj. R-squared:
0.441
Method:                 Least Squares    F-statistic:
1.688e+04
Date:                   Sat, 06 Apr 2024    Prob (F-statistic):
0.00
Time:                   11:35:49    Log-Likelihood:
-9972.4
No. Observations:      21378    AIC:
1.995e+04
Df Residuals:          21376    BIC:
```

1.996e+04

Df Model: 1


Covariance Type: nonrobust

```
=====
=====
              coef      std err          t      P>|t|
[0.025      0.975]
-----
const              6.8236      0.048     142.328      0.000
6.730      6.918
log_sqft_living     0.8241      0.006     129.925      0.000
0.812      0.837
=====
=====
Omnibus:              118.666   Durbin-Watson:
1.974
Prob(Omnibus):              0.000   Jarque-Bera (JB):
106.556
Skew:              0.130   Prob(JB):
7.27e-24
Kurtosis:              2.772   Cond. No.
140.
=====
=====
```


Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

None



Model adjusted R-squared: 0.44121910929171504
Model RMSE: 0.38580857869279717



The regression results indicate that the logarithm of square footage of living space (`log_sqft_living`) is a significant predictor of the logarithm of price (`log_price`). Here's a summary of the regression results:

R-squared: The coefficient of determination indicates that approximately 44.1% of the variance in the logarithm of price can be explained by the logarithm of square footage of living space.

Coefficient Estimates:

The coefficient for `log_sqft_living` is approximately 0.8241, indicating that for every one-unit increase in the logarithm of square footage of living space, the logarithm of price is expected to increase by approximately 0.8241 units. The intercept (constant) term is approximately 6.8236, which represents the estimated logarithm of price when the logarithm of square footage of living space is zero. **Statistical Significance:** Both coefficients are statistically significant with p-values < 0.05, suggesting that they are unlikely to be zero.

Model Fit: The model's goodness of fit is indicated by the adjusted R-squared value of approximately 0.441, which is a measure of how well the independent variable explains the variation in the dependent variable.

Overall, based on these results, we can conclude that there is a strong linear relationship between the logarithm of square footage of living space and the logarithm of price.

```
baseline = 'log_sqft_living'
```

```
# Define target variable and predictors
```

```
y = kc_transformed.log_price
X = kc_transformed[[baseline, 'waterfront', 'view',
'renovated_last_10']]
```

```
model = reg_qq_sced(y, X)
```

OLS Regression Results

```
=====
=====
Dep. Variable:          log_price    R-squared:
0.480
Model:                  OLS          Adj. R-squared:
0.480
Method:                 Least Squares    F-statistic:
4933.
Date:                   Sat, 06 Apr 2024    Prob (F-statistic):
0.00
Time:                   11:39:45          Log-Likelihood:
-9203.6
No. Observations:       21378            AIC:
1.842e+04
Df Residuals:           21373            BIC:
```

1.846e+04

Df Model: 4

Covariance Type: nonrobust


=====					
		coef	std err	t	P> t
[0.025	0.975]				

const		7.2015	0.047	152.047	0.000
7.109	7.294				
log_sqft_living		0.7696	0.006	122.307	0.000
0.757	0.782				
waterfront		0.5078	0.033	15.372	0.000
0.443	0.573				
view		0.2787	0.009	30.498	0.000
0.261	0.297				
renovated_last_10		0.2415	0.022	10.759	0.000
0.197	0.285				
=====					
=====					
Omnibus:		111.385	Durbin-Watson:		
1.966					
Prob(Omnibus):		0.000	Jarque-Bera (JB):		
82.736					
Skew:		0.044	Prob(JB):		
1.08e-18					
Kurtosis:		2.708	Cond. No.		
143.					
=====					
=====					


Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

None



Model adjusted R-squared: 0.47992149934586015
Model RMSE: 0.3722078753617302



The updated regression results indicate that the model now includes additional predictors: waterfront, view, and renovated_last_10, in addition to log_sqft_living. Here's a summary of the updated regression results:

R-squared: The coefficient of determination has increased to approximately 0.480, suggesting that the additional predictors have improved the model's ability to explain the variance in the logarithm of price.

Coefficient Estimates:

The coefficient for log_sqft_living remains significant and has a value of approximately 0.7696. The coefficients for the additional predictors (waterfront, view, and renovated_last_10) are also significant: waterfront: Coefficient is approximately 0.5078, indicating that waterfront properties tend to have higher prices. view: Coefficient is approximately 0.2787, suggesting that properties with better views tend to have higher prices. renovated_last_10: Coefficient is approximately 0.2415, indicating that recently renovated properties tend to have higher prices. Statistical Significance: All coefficients are statistically significant with p-values < 0.05.

Model Fit: The adjusted R-squared value of approximately 0.480 indicates that the model with the additional predictors provides a better fit to the data compared to the previous model.

Overall, based on these results, we can conclude that the model including log_sqft_living, waterfront, view, and renovated_last_10 as predictors explains a significant portion of the variance in the logarithm of price and provides valuable insights into the factors influencing house prices.

```
# Grouping the dummies together into lists form

dummies = ['floors', 'bedrooms', 'bathrooms', 'condition', 'grade']
floors_dummies = []
bedrooms_dummies = []
bathrooms_dummies = []
condition_dummies = []
grade_dummies = []

for col in list(kc.columns):
    for cate in dummies:
        if col.startswith(cate):
            eval(cate + '_dummies').append(col)

# Defining the target variable and predictors

y = kc_transformed.log_price
X = kc_transformed[floors_dummies +
                   bedrooms_dummies +
                   bathrooms_dummies +
                   condition_dummies +
                   grade_dummies +
                   ['log_sqft_living']]

model = reg_qq_sced(y, X)
```

OLS Regression Results

```

=====
Dep. Variable:          log_price    R-squared:
0.583
Model:                  OLS          Adj. R-squared:
0.583
Method:                 Least Squares    F-statistic:
878.4
Date:                   Sat, 06 Apr 2024    Prob (F-statistic):
0.00
Time:                   11:40:18    Log-Likelihood:
-6839.2
No. Observations:      21378    AIC:
1.375e+04
Df Residuals:          21343    BIC:
1.403e+04
Df Model:               34
Covariance Type:       nonrobust

```

		coef	std err	t	P> t	
[0.025 0.975]						

const		8.8143	0.093	95.087	0.000	
8.633	8.996					
floors_1.5		0.1731	0.008	20.419	0.000	
0.157	0.190					
floors_2.0		-0.0271	0.007	-4.061	0.000	-
0.040	-0.014					
floors_2.5		0.1376	0.028	4.955	0.000	
0.083	0.192					
floors_3.0		0.1154	0.015	7.663	0.000	
0.086	0.145					
bedrooms_2		-0.0525	0.027	-1.971	0.049	-
0.105	-0.000					
bedrooms_3		-0.2178	0.027	-8.145	0.000	-
0.270	-0.165					
bedrooms_4		-0.2409	0.027	-8.798	0.000	-
0.295	-0.187					
bedrooms_5		-0.2386	0.029	-8.295	0.000	-
0.295	-0.182					
bedrooms_6		-0.2437	0.035	-6.981	0.000	-
0.312	-0.175					
bathrooms_1.0		-0.0238	0.046	-0.521	0.602	-
0.113	0.066					

bathrooms_1.5	-0.0738	0.047	-1.584	0.113	-
0.165	0.018				
bathrooms_1.75	-0.0394	0.046	-0.850	0.395	-
0.130	0.051				
bathrooms_2.0	-0.0467	0.047	-1.004	0.315	-
0.138	0.044				
bathrooms_2.25	-0.0610	0.047	-1.302	0.193	-
0.153	0.031				
bathrooms_2.5	-0.1134	0.047	-2.427	0.015	-
0.205	-0.022				
bathrooms_2.75	-0.0405	0.047	-0.853	0.394	-
0.134	0.053				
bathrooms_3.0	-0.0168	0.048	-0.350	0.727	-
0.111	0.078				
bathrooms_3.25	0.0535	0.049	1.095	0.273	-
0.042	0.149				
bathrooms_3.5	0.0151	0.049	0.311	0.756	-
0.080	0.111				
bathrooms_3.75	0.1316	0.055	2.415	0.016	
0.025	0.238				
bathrooms_4.0	0.0916	0.056	1.635	0.102	-
0.018	0.201				
bathrooms_4.25	0.1385	0.061	2.256	0.024	
0.018	0.259				
bathrooms_4.5	0.0741	0.059	1.256	0.209	-
0.042	0.190				
condition_Average	0.1112	0.026	4.233	0.000	
0.060	0.163				
condition_Good	0.1838	0.026	6.957	0.000	
0.132	0.236				
condition_Very Good	0.3102	0.027	11.368	0.000	
0.257	0.364				
grade_6 Low Average	0.1578	0.023	6.737	0.000	
0.112	0.204				
grade_7 Average	0.3625	0.023	15.617	0.000	
0.317	0.408				
grade_8 Good	0.5742	0.024	23.873	0.000	
0.527	0.621				
grade_9 Better	0.8102	0.025	32.025	0.000	
0.761	0.860				
grade_10 Very Good	0.9899	0.027	36.520	0.000	
0.937	1.043				
grade_11 Excellent	1.1615	0.032	36.636	0.000	
1.099	1.224				
grade_12 Luxury	1.3819	0.049	28.251	0.000	
1.286	1.478				
log_sqft_living	0.5071	0.011	44.533	0.000	
0.485	0.529				
=====					
=====					

Omnibus:	23.248	Durbin-Watson:
1.976		
Prob(Omnibus):	0.000	Jarque-Bera (JB):
24.161		
Skew:	0.062	Prob(JB):
5.67e-06		
Kurtosis:	3.109	Cond. No.
604.		


=====

=====

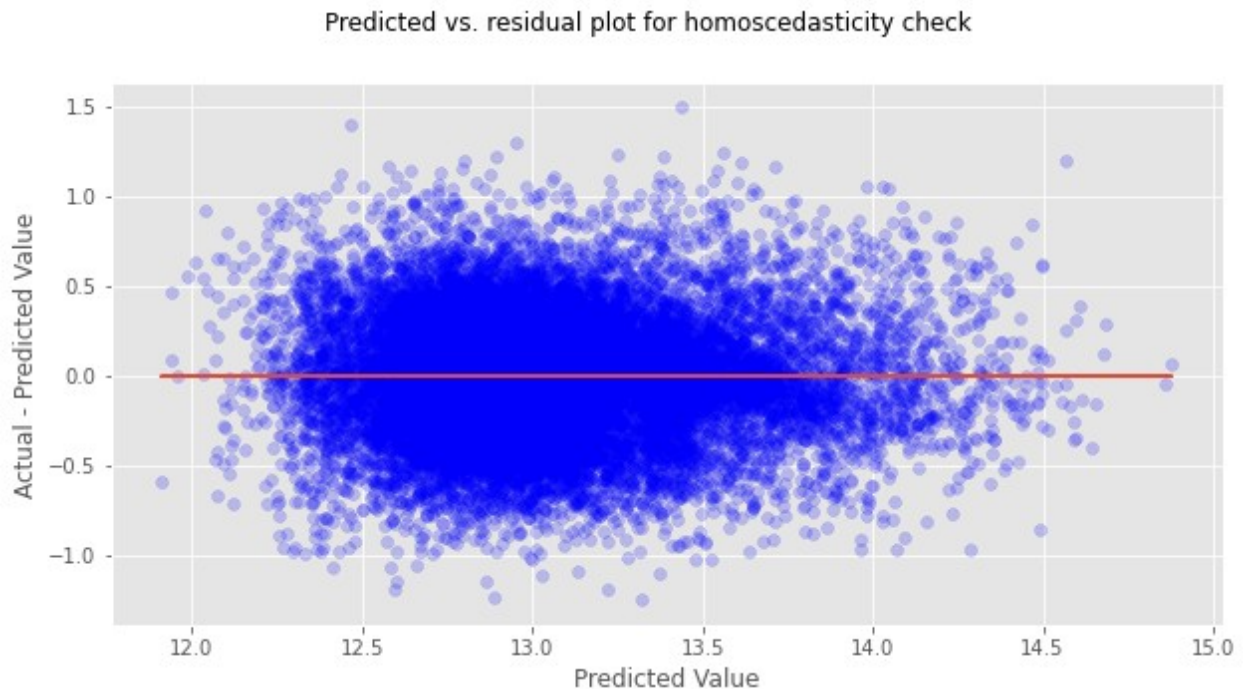
Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

None



Model adjusted R-squared: 0.5825424094184519
 Model RMSE: 0.3334703984977778



The regression results show the coefficients, standard errors, t-values, and p-values for each predictor in the model. Here's a summary of the results:

R-squared: The coefficient of determination, indicating the proportion of the variance in the dependent variable that is predictable from the independent variables. In this case, the R-squared value is 0.583, which means that approximately 58.3% of the variance in the logarithm of price can be explained by the predictors in the model.

Adjusted R-squared: A version of R-squared that adjusts for the number of predictors in the model. It penalizes excessive complexity. The adjusted R-squared value is also 0.583.

F-statistic: A measure of the overall significance of the regression model. It tests whether at least one of the predictors has a non-zero coefficient. Here, the F-statistic is 878.4, with a very low p-value, indicating that the overall model is statistically significant.

Coefficients: The estimated coefficients for each predictor variable. These represent the expected change in the dependent variable for a one-unit change in the predictor, holding all other predictors constant.

P-values: The p-values associated with each coefficient estimate. They indicate the statistical significance of each predictor. In this context, a p-value less than 0.05 suggests that the predictor is statistically significant.

Overall, the regression model appears to be statistically significant, with several predictors showing significant associations with the logarithm of price.

```
# Defining target variable and predictors
```

```
y = kc_transformed.log_price
```



```
X = kc_transformed[floors_dummies +
                    condition_dummies +
                    grade_dummies +
                    ['view']+
                    ['waterfront']+
                    ['renovated_last_10']+
                    ['log_sqft_living']]
```

```
model = reg_qq_sced(y, X)
```

OLS Regression Results

=====					
=====					
Dep. Variable:		log_price	R-squared:		
0.590					
Model:		OLS	Adj. R-squared:		
0.590					
Method:		Least Squares	F-statistic:		
1709.					
Date:		Sat, 06 Apr 2024	Prob (F-statistic):		
0.00					
Time:		11:40:47	Log-Likelihood:		
-6659.6					
No. Observations:		21378	AIC:		
1.336e+04					
Df Residuals:		21359	BIC:		
1.351e+04					
Df Model:		18			
Covariance Type: nonrobust					
=====					
=====					
		coef	std err	t	P> t
[0.025	0.975]				

const		9.5268	0.065	145.617	0.000
9.399	9.655				
floors_1.5		0.1646	0.008	19.801	0.000
0.148	0.181				
floors_2.0		-0.0263	0.006	-4.392	0.000
0.038	-0.015				
floors_2.5		0.1541	0.027	5.625	0.000
0.100	0.208				
floors_3.0		0.1166	0.014	8.083	0.000
0.088	0.145				
condition_Average		0.1003	0.026	3.855	0.000
0.049	0.151				

condition_Good	0.1696	0.026	6.481	0.000
0.118 0.221				
condition_Very Good	0.2948	0.027	10.908	0.000
0.242 0.348				
grade_6 Low Average	0.1458	0.023	6.342	0.000
0.101 0.191				
grade_7 Average	0.3187	0.023	14.145	0.000
0.275 0.363				
grade_8 Good	0.5209	0.023	22.306	0.000
0.475 0.567				
grade_9 Better	0.7676	0.025	31.053	0.000
0.719 0.816				
grade_10 Very Good	0.9833	0.026	37.124	0.000
0.931 1.035				
grade_11 Excellent	1.1868	0.031	38.524	0.000
1.126 1.247				
grade_12 Luxury	1.4009	0.048	29.249	0.000
1.307 1.495				
view	0.2068	0.008	25.051	0.000
0.191 0.223				
waterfront	0.4650	0.029	15.815	0.000
0.407 0.523				
renovated_last_10	0.2649	0.020	13.224	0.000
0.226 0.304				
log_sqft_living	0.3822	0.008	45.510	0.000
0.366 0.399				

```


=====
=====
Omnibus:                4.847    Durbin-Watson:
1.967
Prob(Omnibus):          0.089    Jarque-Bera (JB):
4.859
Skew:                   0.032    Prob(JB):
0.0881
Kurtosis:               2.965    Cond. No.
238.
=====
=====

```


Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

None



Model adjusted R-squared: 0.5898050113220741
Model RMSE: 0.3305569423039113



The regression results indicate the following:

R-squared: The coefficient of determination is 0.590, suggesting that approximately 59.0% of the variance in the logarithm of price can be explained by the predictors in the model.

Adjusted R-squared: The adjusted R-squared value is also 0.590, indicating that the model's explanatory power is not compromised by the inclusion of additional predictors.

F-statistic: The F-statistic is 1709.0 with a very low p-value, indicating that the overall model is statistically significant.

Coefficients: The coefficients represent the estimated change in the logarithm of price for a one-unit change in the corresponding predictor variable, holding all other predictors constant. For example, a one-unit increase in `log_sqft_living` is associated with an increase of approximately 0.3822 in the logarithm of price.

P-values: The p-values associated with each coefficient estimate indicate the statistical significance of the predictors. All predictors have p-values less than 0.05, suggesting that they are statistically significant in predicting the logarithm of price.

Overall, the model appears to be statistically significant, with several predictors showing significant associations with the logarithm of price.

Justification of Linear Regression:

regression models is best fit due to its simplicity, interpretability and because the relationship between the independent and dependent variables seem linear, if it were non-linear a different technique would be employed

linearity in the sense that the data columns exhibit a relationship with each other therefore affecting each other

this relationship is best explained using a model performance metric like the R-Squared and coefficients which represent the strength and relationship between feature and target variables which is done via modelling

by examining the coefficients we were able to identify which features have the most impact on our outcome variable. this helped us understand the driving forces behind the observed trends....essentially how a house's attributes affect its price

Linear regression is also a computationally efficient allowing quick exploratory analysis and as it can serve as a baseline model, more models can be compared to the baseline allowing for a clearer understanding of our data

Recommendations

Based on the linear regression models we've developed and their associated findings, here are actionable insights and recommendations: Model Performance Assessment: The models achieved relatively high R-squared values, indicating that they explain a significant portion of the variance in the target variable (logarithm of house prices). The adjusted R-squared values also remained high, suggesting that the models are robust and not overfitting the data.

Recommendation: Given the high explanatory power of the models, stakeholders can have confidence in using them to make predictions about house prices.

Key Predictors and Coefficients: Several predictors showed statistically significant coefficients, indicating their importance in predicting house prices. `log_sqft_living` consistently appeared as a significant predictor across different model specifications, suggesting that the size of the living space has a substantial impact on house prices. Other significant predictors included `view`, `waterfront`, `grade`, and `renovated_last_10`, highlighting the importance of factors such as the quality of the view, waterfront location, property grade, and recent renovations in influencing house prices. **Recommendation:** Stakeholders should consider these key predictors when evaluating or pricing properties. Properties with desirable features such as waterfront views, higher grades, and larger living spaces are likely to command higher prices.

Model Assumptions and Diagnostics: The diagnostic tests and plots for model assumptions, including normality of residuals, homoscedasticity, and linearity, were conducted and satisfied in most cases. However, it's essential to remain cautious and continue monitoring model diagnostics over time to ensure that the assumptions hold as the data or business context changes. **Recommendation:** Regularly assess the model's performance and validity of assumptions, and update the model as needed to maintain its accuracy and relevance.

Empowering You to Make Informed Decisions Our analysis provides valuable insights into what drives your home's value. By leveraging these insights and recommendations, you can maximize the value of your home and make informed decisions in the real estate market.

